

1. What does `flow_from_directory` give you on the `ImageGenerator`?

- ☐ The ability to easily load images for training
- ☐ The ability to pick the size of training images
- ☐ The ability to automatically label images based on their directory name
- ☒ All of the above

✓ Correct

2. If my Image is sized 150x150, and I pass a 3x3 Convolution over it, what size is the resulting image?

- ☒ 148x148
- ☐ 153x153
- ☐ 450x450
- ☐ 150x150

✓ Correct

3. If my data is sized 150x150, and I use Pooling of size 2x2, what size will the resulting image be?

- ☐ 149x149
- ☐ 148x148
- ☐ 300x300
- ☒ 75x75

✓ Correct

4. If I want to view the history of my training, how can I access it?

- ☐ Pass the parameter `'history=true'` to the `model.fit`
- ☒ Create a variable `'history'` and assign it to the return of `model.fit` or `model.fit_generator`
- ☐ Download the model and inspect it
- ☐ Use a `model.fit_generator`

✓ Correct

5. What's the name of the API that allows you to inspect the impact of convolutions on the images?

- ☐ The model.pools API
- ☐ The model.images API
- ☐ The model.convolution API
- ☒ The model.layers API

✓ Correct

6. When exploring the graphs, the loss levelled out at about .75 after 2 epochs, but the accuracy climbed close to 1.0 after 15 epochs. What's the significance of this?

- ☐ There was no point training after 2 epochs, as we overfit to the validation data
- ☒ There was no point training after 2 epochs, as we overfit to the training data
- ☐ A bigger training set would give us better validation accuracy
- ☐ A bigger validation set would give us better training accuracy

✓ Correct

7. Why is the validation accuracy a better indicator of model performance than training accuracy?

- ☐ It isn't, they're equally valuable
- ☐ There's no relationship between them
- ☒ The validation accuracy is based on images that the model hasn't been trained with, and thus a better indicator of how the model will perform with new images.
- ☐ The validation dataset is smaller, and thus less accurate at measuring accuracy, so its performance isn't as important

✓ Correct

8. Why is overfitting more likely to occur on smaller datasets?

- ☐ Because in a smaller dataset, your validation data is more likely to look like your training data
- ☐ Because there isn't enough data to activate all the convolutions or neurons
- ☐ Because with less data, the training will take place more quickly, and some features may be missed
- ☒ Because there's less likelihood of all possible features being encountered in the training process.